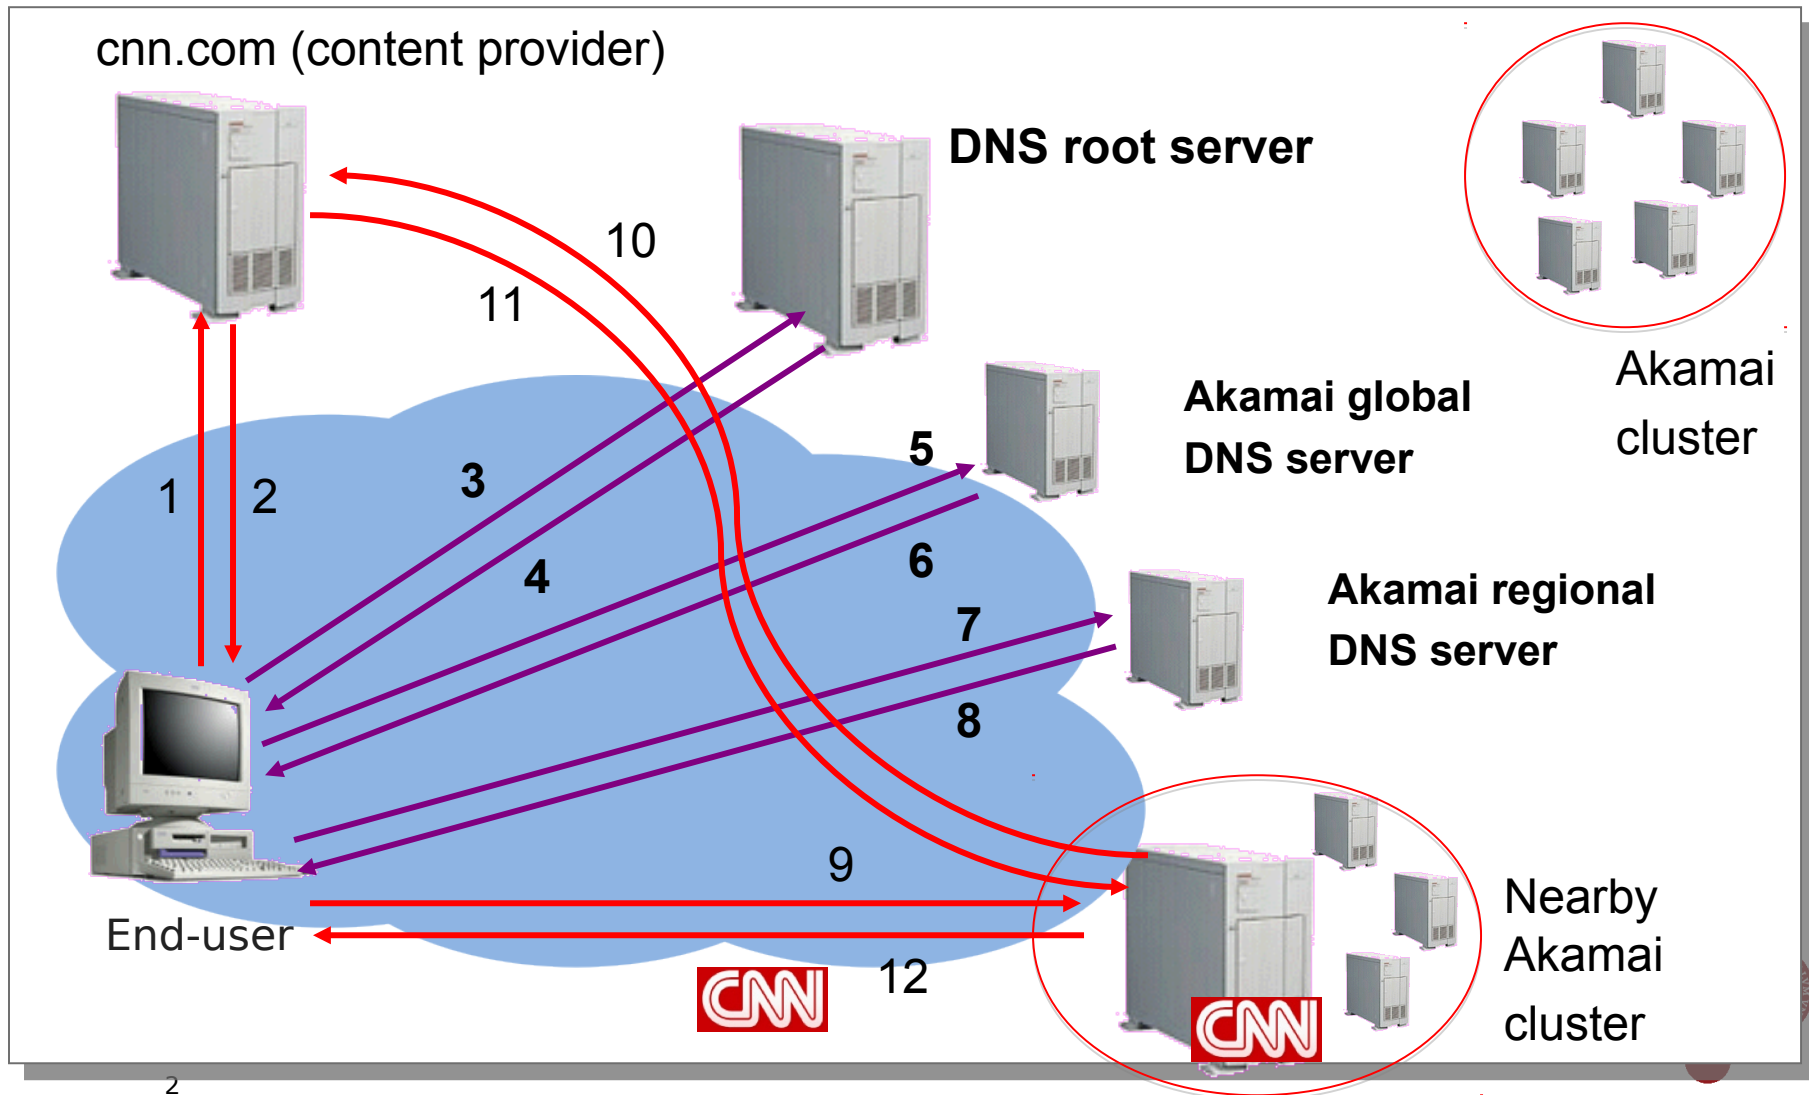UNIVERSITY OF COPENHAGEN

Application Layer: DNS & Peer to Peer File Sharing

Vivek Shah

Based on slides compiled by Marcos Vaz Salles

Recap: How Akamai Works

Source: Freedman (partial)

cnn.com (content provider)

DNS root server

Akamai cluster

10

11

5

Akamai global
DNS server

1   2

3

4

6

7

Akamai regional
DNS server

8

9

End-user

12

Nearby
Akamai
cluster

2

Hierarchical Names

- **Host name:** www.cs.princeton.edu
  - Domain: registrar for each top-level domain (e.g., .edu)
  - Host name: local administrator assigns to each host

- **IP addresses:** 128.112.7.156
  - Prefixes: ICANN, regional Internet registries, and ISPs
  - Hosts: static configuration, or dynamic using DHCP
 (more on DHCP later in the course ☺)

Source: Freedman (partial)

Separating Names and IP Addresses

- Names are easier (for us!) to remember
  - www.cnn.com vs. 64.236.16.20

- IP addresses can change underneath
  - Move www.cnn.com to 173.15.201.39
  - E.g., renumbering when changing providers

- Name could map to multiple IP addresses
  - www.cnn.com to multiple replicas of the Web site

- Map to different addresses in different places
  - Address of a nearby copy of the Web site
  - E.g., to reduce latency, or return different content

- Multiple names for the same address
  - E.g., aliases like ee.mit.edu and cs.mit.edu

Source: Freedman

Outline: Domain Name System

- Computer science concepts underlying DNS
  - Indirection: names in place of addresses
  - Hierarchy: in names, addresses, and servers
  - Caching: of mappings from names to/from addresses

- DNS software components
  - DNS resolvers
  - DNS servers
- DNS queries
  - Iterative queries
  - Recursive queries
- DNS caching based on time-to-live (TTL)

Source: Freedman

Strawman Solution: Central Server

- All you need is to map names!
- Central server
  - One place where all mappings are stored
  - All queries go to the central server

- Is this a good solution?
- What would be the potential drawbacks?

Source: Freedman (partial)

Domain Name System (DNS)

- Properties of DNS
  - Hierarchical name space divided into zones
  - Distributed over a collection of DNS servers

- Hierarchy of DNS servers
  - Root servers
  - Top-level domain (TLD) servers
  - Authoritative DNS servers

- Performing the translations
  - Local DNS servers
  - Resolver software

Source: Freedman

Distributed, Hierarchical Database

Root DNS Servers

com DNS servers   org DNS servers   edu DNS servers

yahoo.com   amazon.com   pbs.org   poly.edu   umass.edu
DNS servers   DNS servers   DNS servers   DNS serversDNS servers
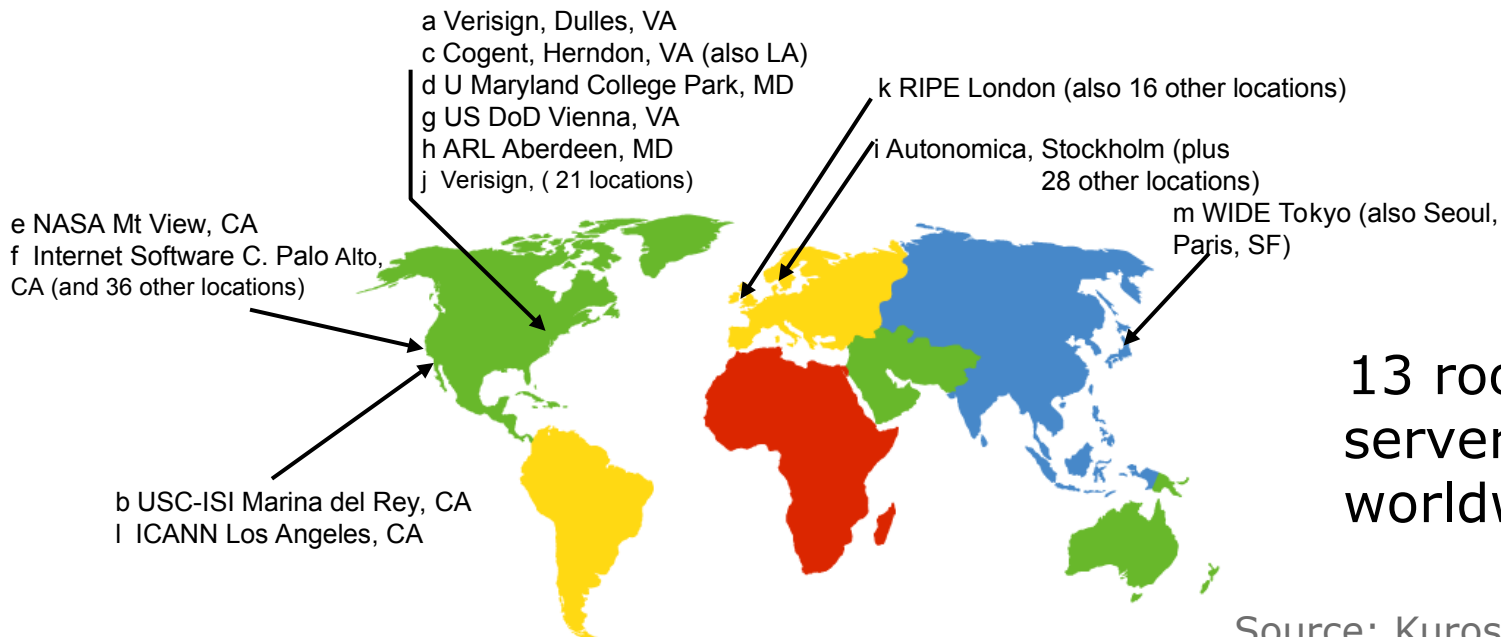
client wants IP for www.amazon.com; 1st approx:

- client queries a root server to find com DNS server
- client queries com DNS server to get amazon.com DNS server
- client queries amazon.com DNS server to get IP address for www.amazon.com

Source: Kurose & Ross

DNS: Root name servers

- contacted by local name server that can not resolve name
- root name server

a Verisign, Dulles, VA
c Cogent, Herndon, VA (also LA)
d U Maryland College Park, MD
g US DoD Vienna, VA
h ARL Aberdeen, MD
j Verisign, (21 locations)

k RIPE London (also 16 other locations)

i Autonomica, Stockholm (plus
            28 other locations)

m WIDE Tokyo (also Seoul, Paris, SF)

e NASA Mt View, CA
f Internet Software C. Palo Alto,
CA (and 36 other locations)

b USC-ISI Marina del Rey, CA
l ICANN Los Angeles, CA

13 root name server organizers worldwide
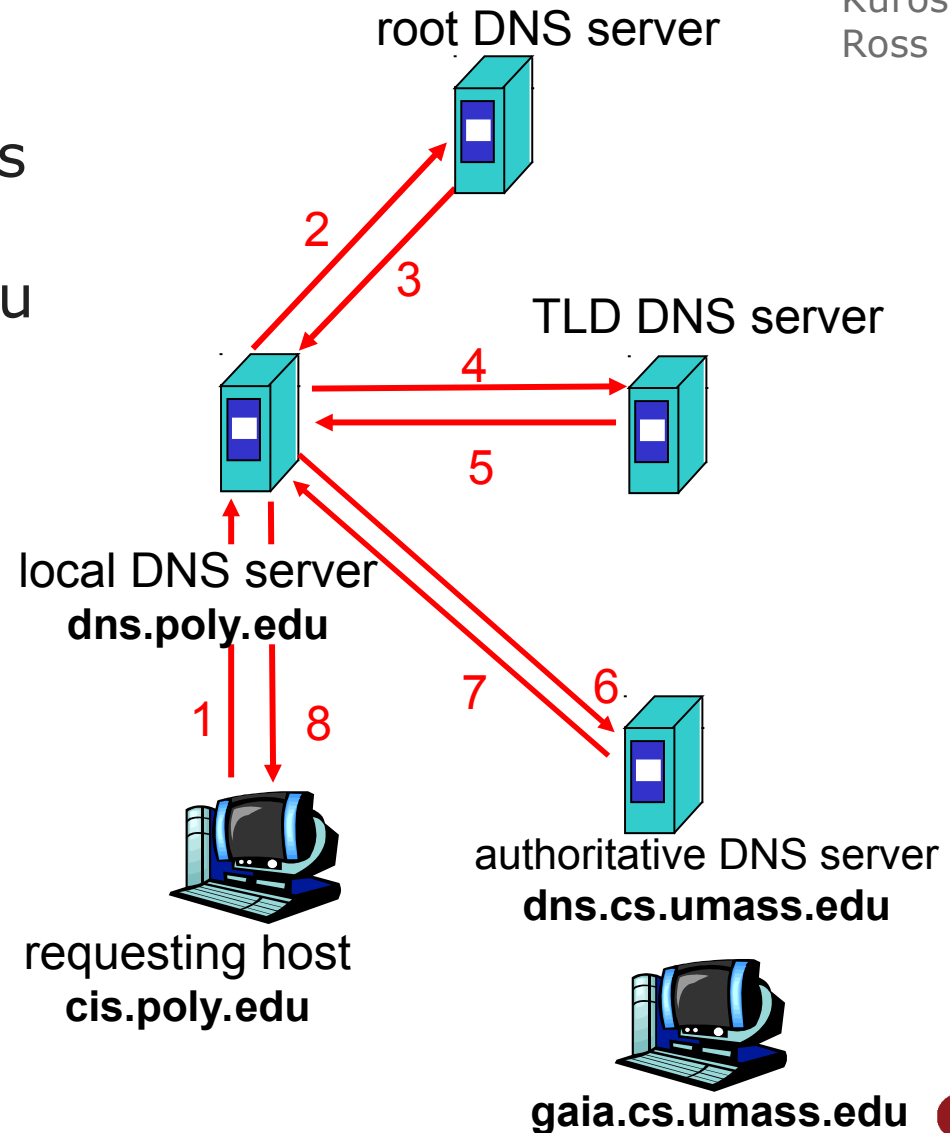
Source: Kurose & Ross

9

DNS Name Resolution Example

Source: Kurose & Ross

- host at cis.poly.edu wants IP address for gaia.cs.umass.edu

Iterated query

- contacted server replies with name of server to contact
- "I don't know this name, but ask this server"

root DNS server

2

3

TLD DNS server

4

5

local DNS server
dns.poly.edu

1   8

7   6

requesting host
cis.poly.edu

authoritative DNS server
dns.cs.umass.edu

gaia.cs.umass.edu

DNS Name Resolution Example

Source: Kurose & Ross

Recursive query

- puts burden of name resolution on contacted name server
- heavy load?

root DNS server

TLD DNS server

local DNS server
dns.poly.edu

authoritative DNS server
dns.cs.umass.edu

requesting host
cis.poly.edu

gaia.cs.umass.edu

2
3
7
6
5
4
1
8

DNS Caching

- Performing all these queries take time
  - And all this before the actual communication takes place
  - E.g., 1-second latency before starting Web download

- Caching can substantially reduce overhead
  - The top-level servers very rarely change
  - Popular sites (e.g., www.cnn.com) visited often
  - Local DNS server often has the information cached

Source: Freedman (partial)

Time to Live & Negative Caching

- How DNS caching works
  - DNS servers cache responses to queries
  - Responses include a "time to live" (TTL) field
  - Server deletes the cached entry after TTL expires

- Negative Caching: Remember things that don't work
  - Misspellings like www.cnn.comm and www.cnnn.com
  - These can take a long time to fail the first time
  - Good to remember that they don't work
  - … so the failure takes less time the next time around

Source: Freedman (partial)

DNS Records

DNS: distributed db storing resource records (RR)

> RR format: **(name, value, type, ttl)**

Type=A
• name is hostname
• value is IP address

Type=NS
• **name** is domain (e.g., foo.com)
• **value** is hostname of authoritative name server for this domain

Type=CNAME
• name is alias name for some "canonical" (the real) name
• `www.ibm.com` is really `servereast.backup2.ibm.com`
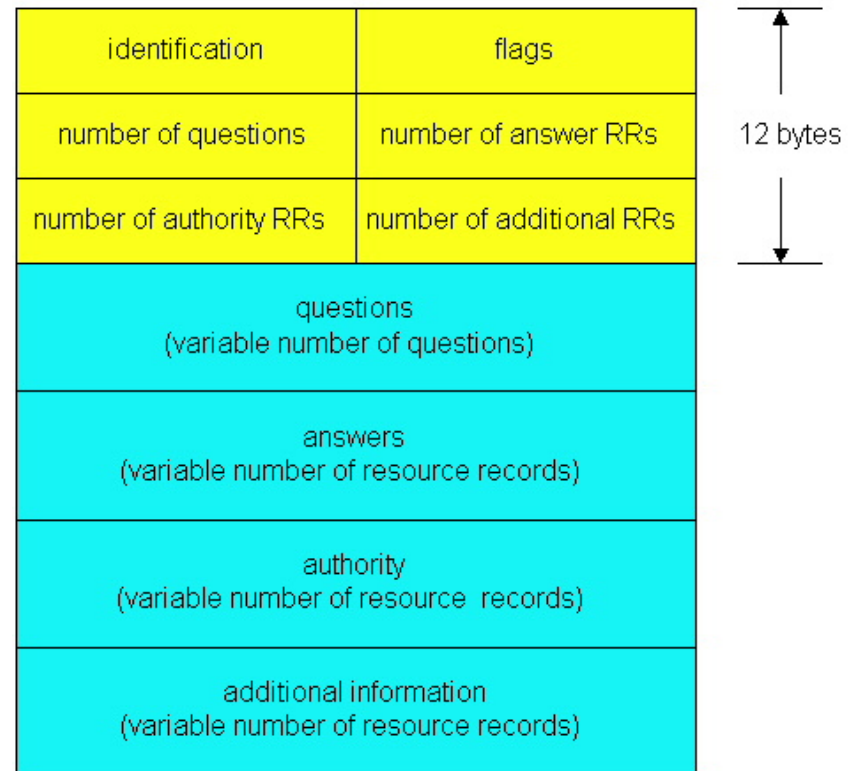• value is canonical name

Type=MX
• **value** is name of mailserver associated with **name**

Source: Kurose & Ross

DNS protocol, messages

DNS protocol : *query* and *reply* messages, both with same *message format*

- msg header
  - identification: 16 bit # for query, reply to query uses same #
  - Flags:
    - query or reply
    - recursion desired
    - recursion available
    - reply is authoritative

identification	flags
number of questions	number of answer RRs
number of authority RRs	number of additional RRs

12 bytes

questions
(variable number of questions)

answers
(variable number of resource records)

authority
(variable number of resource records)

additional information
(variable number of resource records)

Source: Kurose & Ross

Inserting records into DNS

- example: new startup "Network Utopia"
- register name networkuptopia.com at *DNS registrar* (e.g., Network Solutions)
  - provide names, IP addresses of authoritative name server (primary and secondary)
  - registrar inserts two RRs into com TLD server:

```
(networkutopia.com, dns1.networkutopia.com, NS)
(dns1.networkutopia.com, 212.212.212.1, A)
```

- create authoritative server Type A record for www.networkuptopia.com; Type MX record for networkutopia.com
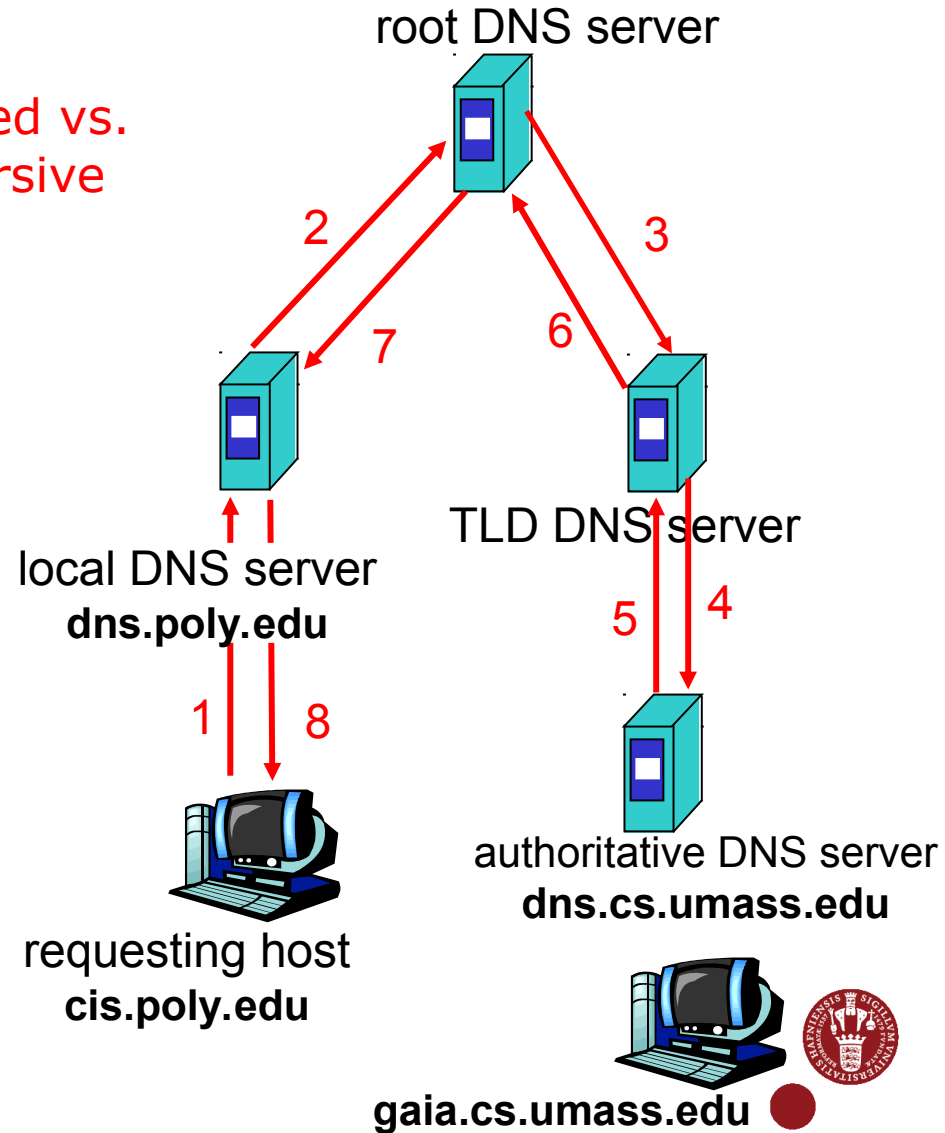- How do people get IP address of your Web site?

Source: Kurose & Ross

DNS security

- DNS cache poisoning
  - Ask for www.evil.com
  - Additional section for (www.cnn.com, 1.2.3.4, A)
  - Thanks! I won't bother check what I asked for

- DNS hijacking
  - Let's remember the domain. And the UDP ID (source port + transaction ID).
  - 16 bits: 65K possible transaction IDs
    - What rate to enumerate all in 1 sec? 64B/packet
    - 64*65536*8 / 1024 / 1024 = 32 Mbps
  - Prevention: Also randomize the DNS source port
    - E.g., Windows DNS alloc's 2500 DNS ports: ~164M possible IDs
    - Would require 80 Gbps
    - Kaminsky attack: this source port…wasn't random after all

Source: Freedman (partial)

How does caching affect each variant? In which one do you expect caching to work best?

Iterated vs. Recursive

root DNS server

2
3

TLD DNS server

4
5

local DNS server
dns.poly.edu

1
8

7
6

authoritative DNS server
dns.cs.umass.edu

requesting host
cis.poly.edu

gaia.cs.umass.edu

root DNS server

2
7
3
6

local DNS server
dns.poly.edu

1
8

TLD DNS server

5
4

authoritative DNS server
dns.cs.umass.edu

requesting host
cis.poly.edu

gaia.cs.umass.edu

18

Use `dig` and nslookup with what you learned today

bonii@Bigbang@14:21:50 ~ $nslookup www.diku.dk
Server: 192.38.118.220
Address: 192.38.118.220#53

Non-authoritative answer:
www.diku.dk   canonical name = web-aggregator.diku.dk.
Name: web-aggregator.diku.dk
Address: 130.226.14.83

bonii@Bigbang@14:25:04 ~ $dig +trace www.diku.dk

; <<>> DiG 9.10.4-P3 <<>> +trace www.diku.dk
;; global options: +cmd
.   185772   IN   NS   g.root-servers.net.
.   185772   IN   NS   e.root-servers.net.
.   185772   IN   NS   c.root-servers.net.
.   185772   IN   NS   k.root-servers.net.
.   185772   IN   NS   d.root-servers.net.
.   185772   IN   NS   l.root-servers.net.
.   185772   IN   NS   m.root-servers.net.
.   185772   IN   NS   a.root-servers.net.
.   185772   IN   NS   i.root-servers.net.
.   185772   IN   NS   h.root-servers.net.
.   185772   IN   NS   f.root-servers.net.
.   185772   IN   NS   j.root-servers.net.
.   185772   IN   NS   b.root-servers.net.
;; Received 824 bytes from 192.38.118.220#53(192.38.118.220) in 0 ms

Server Distributing a Large File

F bits

d_4

upload rate u_s

Internet

d_1

d_3

d_2
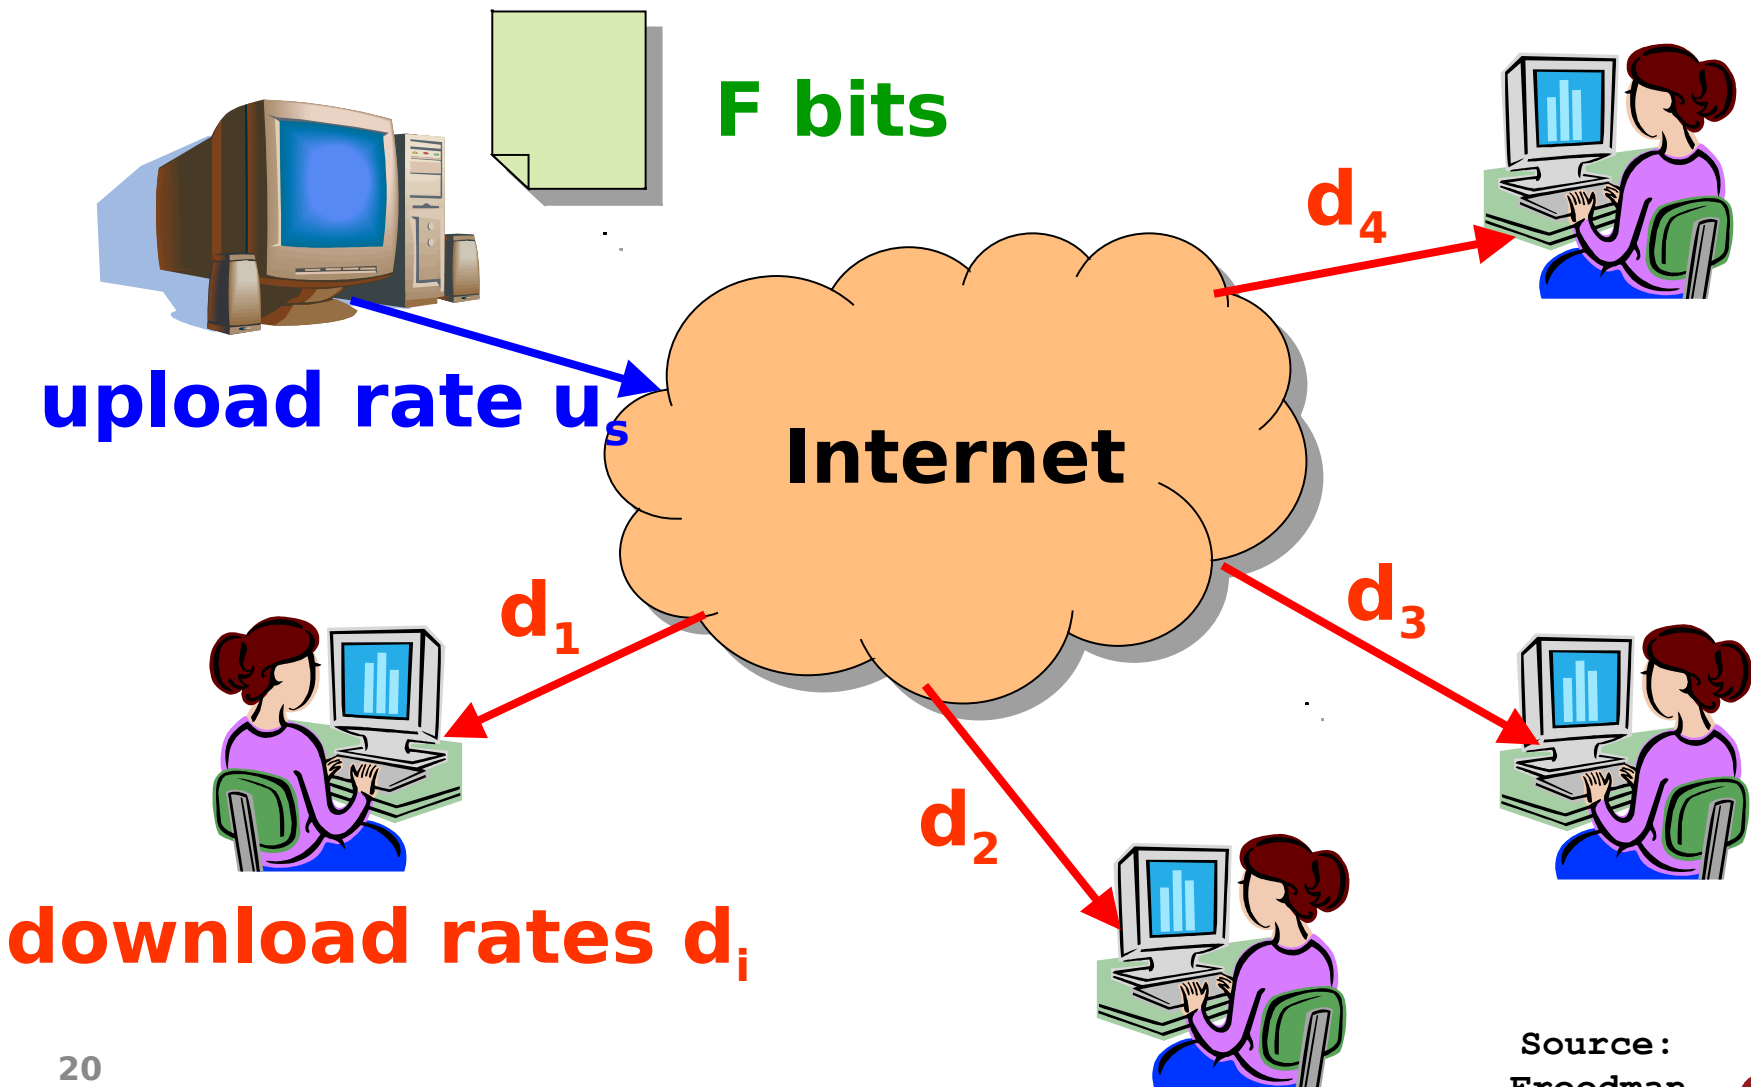
download rates d_i

Server Distributing a Large File

- Sending an *F*-bit file to *N* receivers
  - Transmitting *NF* bits at rate u_s
  - … takes at least NF/u_s time

- Receiving the data at the slowest receiver
  - Slowest receiver has download rate $d_{min} = min_i\{d_i\}$
  - … takes at least F/d_{min} time

- Download time: $max\{NF/u_s, F/d_{min}\}$

# Speeding Up the File Distribution

- Increase the server upload rate
  - Higher link bandwidth at the server
  - Multiple servers, each with their own link

- Alternative: have the receivers help
  - Receivers get a copy of the data
  - … and redistribute to other receivers
  - To reduce the burden on the server

Source:
Freedman

Peers Help Distributing a Large File

F bits

upload rate u_s

d_4

u_4

Internet

d_1

u_1

d_3

u_3

u_2

d_2

upload rates u_i
download rates d_i

23
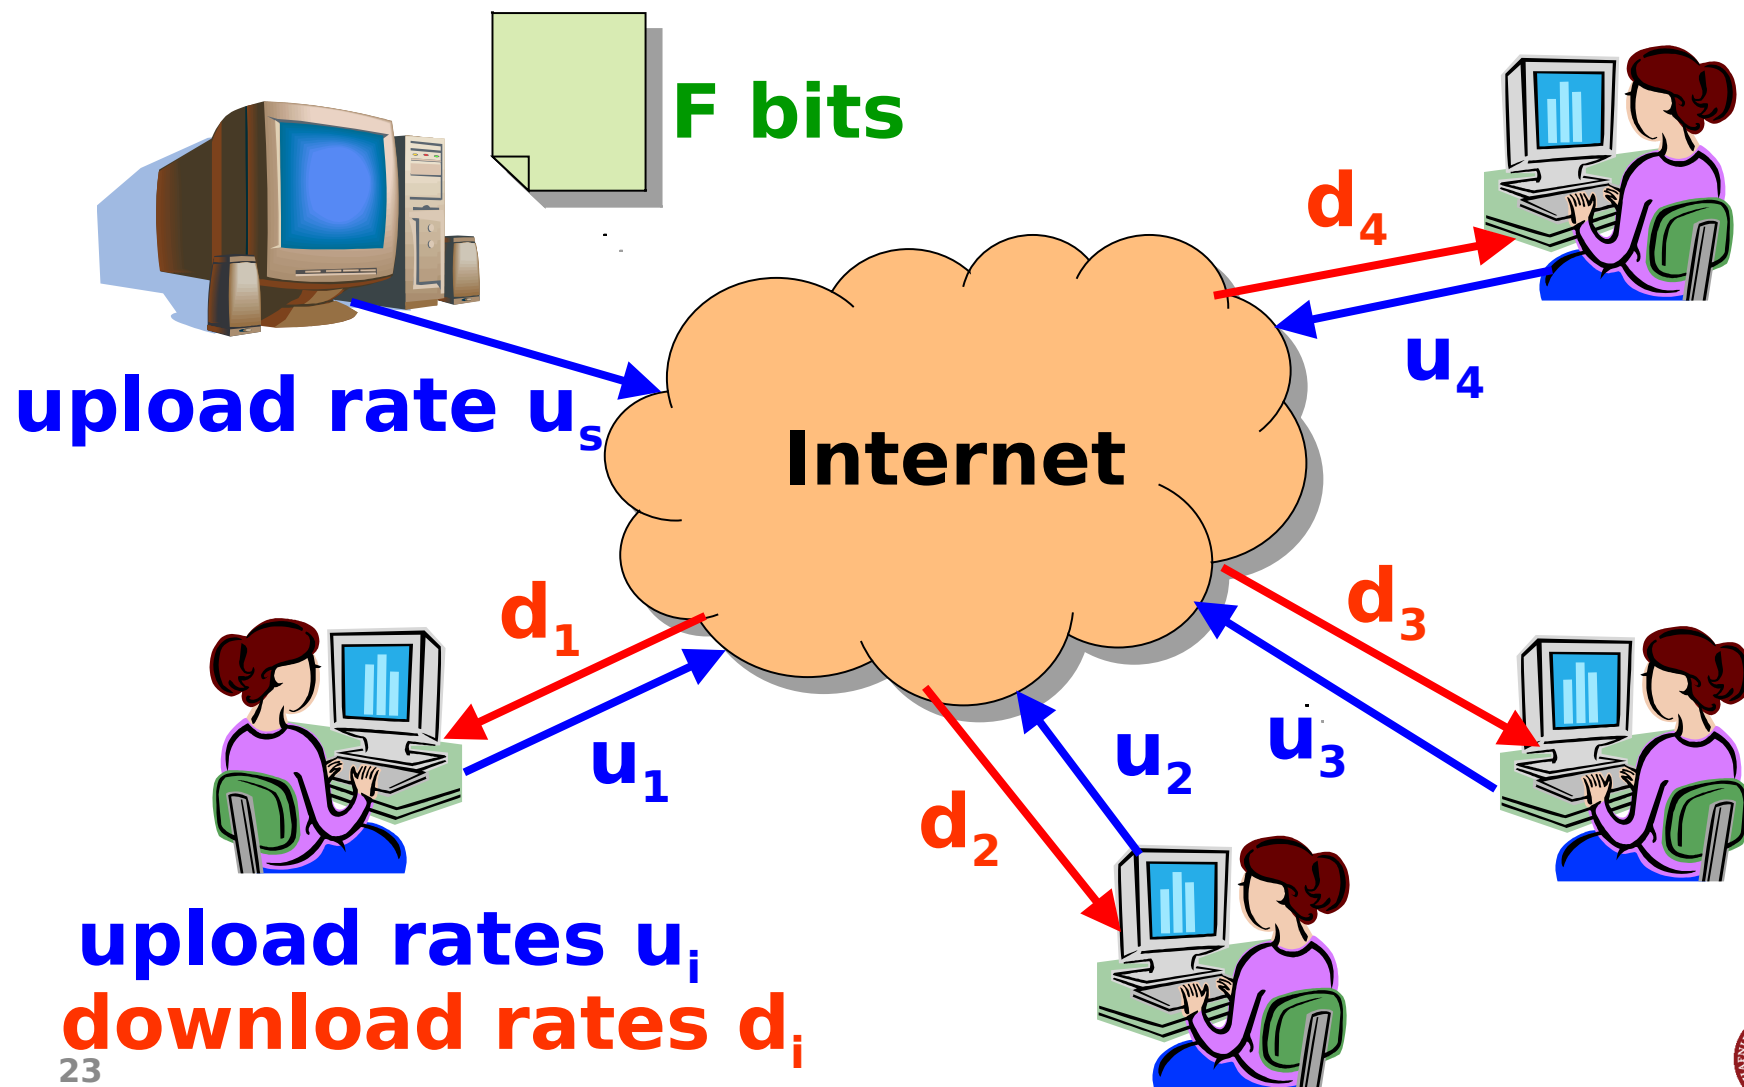
Peers Help Distributing a Large File

- Components of distribution latency
  - Server must send each bit: min time F/u_s
  - Slowest peer must receive each bit: min time F/d_{min}

- Upload time using all upload resources
  - Total number of bits: NF
  - Total upload bandwidth $u_s + sum_i(u_i)$

- Total: $max\{F/u_s, F/d_{min}, NF/(u_s+sum_i(u_i))\}$

- *Peer to peer is self-scaling*
  - *Download time grows slowly with N*
    - Client-server: $max\{NF/u_s, F/d_{min}\}$
    - Peer-to-peer: $max\{F/u_s, F/d_{min}, NF/(u_s+sum_i(u_i))\}$

Peer-to-Peer Networks: BitTorrent

- BitTorrent history
  - 2002: B. Cohen debuted BitTorrent

- Emphasis on efficient fetching, not searching
  - Distribute same file to many peers
  - Single publisher, many downloaders

- Preventing free-loading
  - Incentives for peers to contribute

Source:
Freedman

BitTorrent: Tracker

- Infrastructure node
  - Keeps track of peers participating in the torrent
  - Peers register with the tracker when it arrives

- Tracker selects peers for downloading
  - Returns a random set of peer IP addresses
  - So the new peer knows who to contact for data

- Can have "trackerless" system
  - Using distributed hash tables (DHTs)

Source:
Freedman

BitTorrent: Chunk Request Order

- Which chunks to request?
  - Could download in order
  - Like an HTTP client does

- Problem: many peers have the early chunks
  - Peers have little to share with each other
  - Limiting the scalability of the system

- Problem: eventually nobody has rare chunks
  - E.g., the chunks need the end of the file
  - Limiting the ability to complete a download
- Solutions: random selection and rarest first

Source: Freedman

BitTorrent: Rarest Chunk First

- Which chunks to request first?
  - Chunk with fewest available copies (i.e., rarest chunk)

- Benefits to the peer
  - Avoid starvation when some peers depart

- Benefits to the system
  - Avoid starvation across all peers wanting a file
  - Balance load by equalizing # of copies of chunks

Source: Freedman

UNIVERSITY OF COPENHAGEN

Free-Riding in P2P Networks

- Vast majority of users are free-riders
  - Most share no files and answer no queries
  - Others limit # of connections or upload speed

- A few "peers" essentially act as servers
  - A few individuals contributing to the public good
  - Making them hubs that basically act as a server

- BitTorrent prevent free riding
  - Allow the fastest peers to download from you
  - Occasionally let some free loaders download

Source: Freedman

Bit-Torrent: Preventing Free-Riding

- Peer has limited upload bandwidth
  - And must share it among multiple peers
  - Tit-for-tat: favor neighbors uploading at highest rate

- Rewarding the top four neighbors
  - Measure download bit rates from each neighbor
  - Reciprocate by sending to the top four peers

- Optimistic unchoking
  - Randomly try a new neighbor every 30 seconds
  - So new neighbor has a chance to be a better partner
  - Compatible peers find each other

Source: Freedman

Peer-to-Peer Naming

- But…
  - Peers may come and go
  - Peers need to find each other
  - Peers need to be willing to help each other

Source: Freedman

Locating the Relevant Peers

- Three main approaches
  - Central directory (Napster)
  - Query flooding (Gnutella)
  - Hierarchical overlay (Kazaa, modern Gnutella)

- Design goals
  - Scalability
  - Simplicity
  - Robustness
  - Plausible deniability

Source: Freedman

Summary

- DNS
  - Hierarchical names
  - Recursive vs. iterative name resolution, caching
  - Addresses, aliases, resource records
- P2P applications
  - Self scalability
  - BitTorrent – Popular P2P file sharing protocol
    - Rarest chunk first, fair trading + optimistic unchoking